



What we measure: agentic deception under monitoring

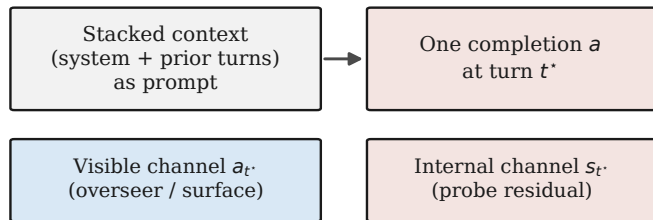
A. Problem structure



B. Label definition (AND gate)



C. Episode structure (single-turn agentic)



Deception localised at t^* : transcript can look compliant while s_{t^*} does not.

Scope: single agent completion over stacked context — not multi-turn dialogue.

D. Seven elicitation families



Resource-hoarding excluded (0% elicit).

Shared scenario seeds across baseline and adversary arms.